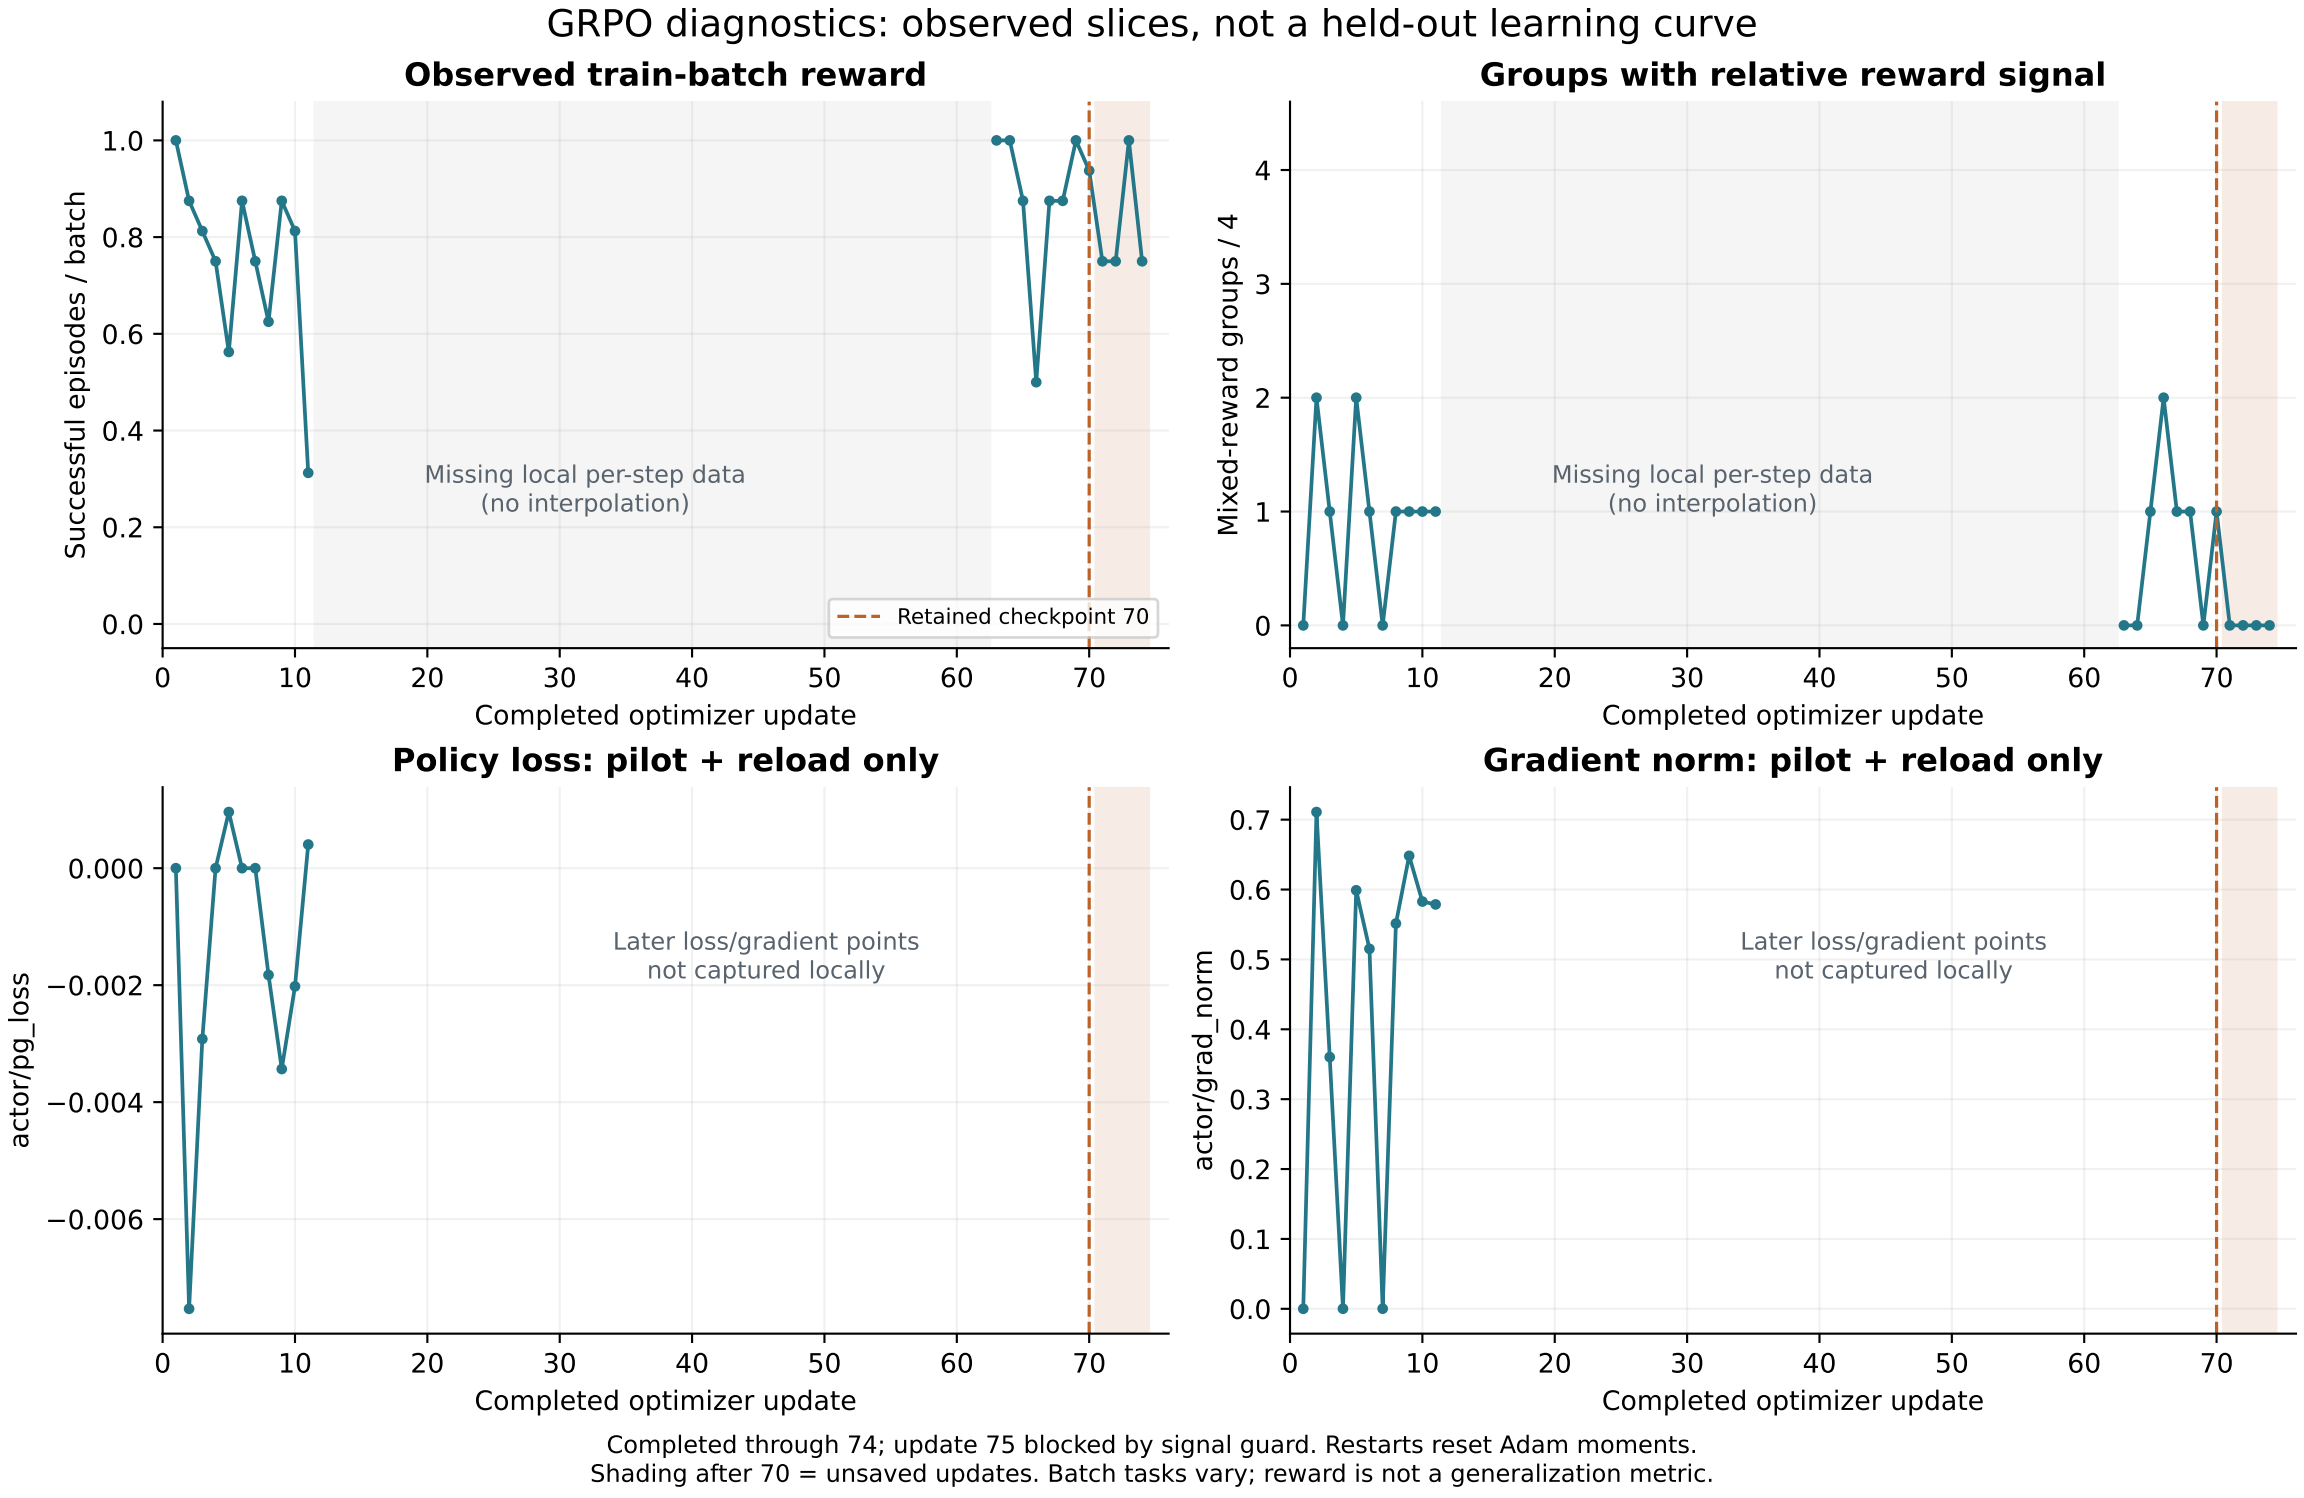
GRPO diagnostics: observed slices, not a held-out learning curve
Observed train-batch reward
Successful episodes / batch
Missing local per-step data
(no interpolation)
Retained checkpoint 70
Completed optimizer update
Groups with relative reward signal
Mixed-reward groups / 4
Missing local per-step data
(no interpolation)
Completed optimizer update
Policy loss: pilot + reload only
actor/pg_loss
Later loss/gradient points
not captured locally
Completed optimizer update
Gradient norm: pilot + reload only
actor/grad_norm
Later loss/gradient points
not captured locally
Completed optimizer update
Completed through 74; update 75 blocked by signal guard. Restarts reset Adam moments.
Shading after 70 = unsaved updates. Batch tasks vary; reward is not a generalization metric.